Regime-Switching Equity Portfolio with Adaptive Allocation

A Deep Reinforcement Learning Approach to Intraday Trading in Indian Markets

Quantitative Finance Project
Advanced Quantitative Analytics

January 10, 2026

Abstract

This report presents a comprehensive regime-aware portfolio strategy that combines Hidden Markov Models (HMM), EGARCH volatility modeling, CUSUM and Deep Reinforcement Learning (DRL) for adaptive equity allocation in Indian markets. Financial markets are fundamentally non-stationary—strategies optimized for low-volatility bull markets often fail catastrophically during high-volatility crashes. Our solution addresses this through a three-tiered "Hybrid Intelligence" architecture: (1) a multi-factor stock selection engine using momentum, volatility, and liquidity scoring across Nifty 50/Next 50 constituents, (2) a walk-forward regime detection framework leveraging HMM and XGBoost for market state classification, and (3) a Proximal Policy Optimization (PPO) agent trained on 200+ high-frequency features for intraday execution. The strategy employs risk-parity weighting and dynamic position sizing to minimize drawdowns while preserving compounding returns.

Keywords: Regime Detection, Hidden Markov Models, Reinforcement Learning, PPO, Algorithmic Trading, Risk Parity, Indian Equities

Contents

1 Introduction

1.1 Problem Statement

Modern portfolio management faces the dual challenge of adapting to varying market conditions while maintaining robust risk controls. Financial markets exhibit fundamental non-stationarity—the statistical properties of returns, volatility, and correlations shift dramatically across different market regimes. A strategy that generates alpha during low-volatility bull markets often experiences catastrophic losses during high-volatility crashes or prolonged sideways consolidation periods.

Traditional static allocation strategies suffer from several critical flaws:

- **Regime Blindness:** Fixed rules (e.g., "Buy when RSI < 30") fail during regime transitions when underlying return distributions change
- **Look-Ahead Bias:** Many backtests inadvertently use future information, inflating performance metrics
- **Over-Fitting:** Complex strategies optimized on historical data often deteriorate in live trading
- **Concentration Risk:** Single-factor strategies expose portfolios to systematic crashes

This project addresses these challenges by developing an intelligent trading system that:

- Detects market regimes (bull, bear, sideways) using probabilistic models without look-ahead bias
- Dynamically selects high-quality stocks from the Nifty 50/Next 50 universe using multi-factor ranking
- Executes intraday trades using deep reinforcement learning with sophisticated reward engineering
- Implements risk-parity allocation to prevent sector concentration
- Minimizes drawdowns through adaptive stop-loss mechanisms and trend structure validation

1.2 Key Innovations

- **Hybrid Intelligence Architecture:** Integration of probabilistic regime detection (HMM) with adaptive execution (PPO)
- **Walk-Forward Validation:** Rolling window training on recent 50% of days, testing on next 50% days to prevent overfitting
- **Sortino-Style Reward Function:** Asymmetric penalty for downside volatility while rewarding upside momentum
- **Trend Structure Veto:** XGBoost classifier acts as secondary filter, vetoing counter-trend signals during structural downtrends
- **200+ Feature Engineering Pipeline:** Multi-timeframe indicators with strict causality constraints

2 Methodology

2.1 System Architecture

The system consists of three interconnected modules operating in a hierarchical pipeline. The market data layer feeds into parallel stock selection and regime detection modules, which then inform the reinforcement learning trading agent for optimal execution decisions.

2.2 Stock Selection Engine

2.2.1 Universe Definition

The investable universe comprises 99 stocks from Nifty 50 and Nifty Next 50 indices, ensuring:

- Adequate liquidity (average daily volume > 1M shares)
- Minimum 30-day trading history
- Availability of complete OHLCV data

2.2.2 Multi-Factor Scoring

Each trading day, stocks are ranked using a composite score:

$$S_i = \sum_{k=1}^{6} w_k \cdot \frac{F_{ik} - \mu_k}{\sigma_k} \tag{1}$$

where:

- F_{ik}: Raw value of factor k for stock i
- w_k: Factor weight (calibrated via backtesting)
- μ_k, σ_k: Cross-sectional mean and standard deviation

Factor Definitions:

1. **Long-term Momentum** ($w = 0.30$): 30-day price change

$$M_{\text{long}} = \frac{P_t - P_{t-30}}{P_{t-30}} \tag{2}$$

2. **Short-term Momentum** ($w = 0.20$): 20-day price change

$$M_{20} = \frac{P_t - P_{t-20}}{P_{t-20}} \tag{3}$$

3. **Realized Volatility** ($w = 0.30$): Annualized log-return std

$$\sigma_r = \text{std}(\ln(P_t/P_{t-1})) \cdot \sqrt{252} \tag{4}$$

4. **Average True Range** ($w = 0.05$): 14-day ATR

$$ATR = \frac{1}{14}\sum_{i=1}^{14} TR_i, \quad TR = \max(H - L, |H - C_{\text{prev}}|, |L - C_{\text{prev}}|) \tag{5}$$

5. **Liquidity** ($w = 0.10$): Mean dollar volume

$$L = \frac{1}{30} \sum_{i=1}^{30} P_i \cdot V_i \tag{6}$$

6. **Market Beta** ($w = 0.05$): Covariance with Nifty 50

$$\beta = \frac{\mathrm{Cov}(R_{\mathrm{stock}}, R_{\mathrm{market}})}{\mathrm{Var}(R_{\mathrm{market}})} \tag{7}$$

Top $K = 10$ stocks are selected daily for trading.

2.3 Regime Detection Framework

2.3.1 Hidden Markov Model

A 3-state Gaussian HMM captures latent market states:

$$P(O_t|S_t = s) = \mathcal{N}(\mu_s, \Sigma_s) \tag{8}$$

Observations: $O_t = [\mathrm{LogRet}_t, \mathrm{GK\_Vol}_t, \mathrm{Vol\_Z}_t]$
States: $S \in \{\text{Low Vol}, \text{Normal}, \text{High Vol}\}$
Features:

- Garman-Klass Volatility: $\mathrm{GK} = \sqrt{0.5(\ln H/L)^2 - (2\ln 2 - 1)(\ln C/O)^2}$
- Volume Z-score: $\mathrm{Vol\_Z} = (V_t - \bar{V})/\sigma_V$

2.3.2 XGBoost Classifier

Trained on 500,000+ samples with SMOTE oversampling to predict future regime (5-minute horizon):

Target Labels:

$$\mathrm{Regime}_t = \begin{cases} \text{Bull} & \text{if } R_{t+5} > 0.03\% \\ \text{Bear} & \text{if } R_{t+5} < -0.03\% \\ \text{Sideways} & \text{otherwise} \end{cases} \tag{9}$$

Feature Set (50+ indicators):

- HMM states and transition probabilities
- CUSUM detectors on returns and volatility
- Multi-timeframe momentum (5m, 15m, 30m, 60m)
- Technical indicators: RSI, MACD, CCI, CMO
- EGARCH conditional volatility
- Exponential moving average crosses (EMA-9, EMA-21)

2.4 Reinforcement Learning Trading Agent

2.4.1 Markov Decision Process Formulation

State Space ($\mathbb{R}^{200+}$):

- **Lookback Features** (10-step window): RSI, CCI, CMO, ATR, Standard Deviation, KAMA & ER.
- **Heikin-Ashi Smoothing** (10 iterations): Trend detection with noise reduction
- **Johnny Ribbon Regimes**: 7 EMA periods with reference MA for multi-timeframe confluence
- **Agent State**: Current position $\{-1, 0, 1\}$, unrealized PnL
- **Regime Prediction**: XGBoost output $\{0, 1, 2\}$ (Bull/Sideways/Bear)
- **Volume Microstructure**: Relative volume, VWAP deviation, Hurst exponent

Action Space: $\mathcal{A} = \{\text{Hold/Exit}, \text{Buy}, \text{Sell}\}$

Reward Function:

Standard RL environments often fail in finance because they penalize *all* volatility equally (Sharpe ratio approach). However, traders desire upside volatility (breakouts) while avoiding downside crashes. We designed a custom asymmetric reward function that separates "Good Volatility" from "Bad Volatility":

$$R_t = \begin{cases} \alpha \cdot \text{PnL}_{\text{bps}} & \text{if PnL} > 0 \text{ (Linear reward)} \\ \lambda \cdot (\text{PnL}_{\text{bps}})^2 \cdot \text{sign(PnL)} & \text{if PnL} < 0 \text{ (Quadratic penalty)} \end{cases} + R_{\text{structural}} \tag{10}$$

Where the structural reward components are:

$$R_{\text{structural}} = \begin{cases} -100 & \text{if stop-loss triggered (catastrophic loss)} \\ -10 & \text{if trailing stop hit (momentum reversal)} \\ -10 & \text{if forced EOD close (overnight risk)} \\ -2 & \text{if opportunity missed (inaction during volatility)} \\ +0.1 & \text{if correctly waiting (avoiding false signals)} \end{cases} \tag{11}$$

Parameters:

- $\alpha = 0.01$: Linear scaling for profits
- $\lambda = 0.04$: Quadratic penalty coefficient (4x base multiplier)
- Loss squaring: A -1% crash hurts 100× more than a -0.1% slip

Impact: By squaring negative returns, the agent learns to exit positions immediately upon trend reversal, effectively developing autonomous stop-loss behavior without explicit programming. This successfully trained the agent to be risk-averse regarding downside volatility while aggressively capturing upside momentum.

Table 1: PPO Hyperparameters

Parameter	Value
Policy Network	MLP [256, 256]
Value Network	MLP [256, 256]
Learning Rate	Linear decay: $3 \times 10^{-3} \rightarrow 1 \times 10^{-4}$
Batch Size	4096
Steps per Env	1024
GAE Lambda	0.90
Discount Factor	0.99
Clip Range	0.2
Entropy Coefficient	0.001
Parallel Envs	8
Total Timesteps	1,000,000

2.4.2 PPO Training Configuration

2.4.3 Risk Management

Training Mode:

- Stop Loss: -0.05% (-5 bps)
- Trailing Stop: 0.05% (5 bps)
- Take Profit: 100% (disabled for exploration)

Testing Mode:

- Stop Loss: -0.04% (-4 bps)
- Trailing Stop: 0.04% (4 bps)
- Forced EOD closing to avoid overnight risk

3 Results

3.1 Performance Overview

The trading system was rigorously backtested on two major constituents of the Nifty 50 index: **HDFCBANK** and **TCS**. The strategy demonstrated exceptional stability across different market regimes, achieving high risk-adjusted returns with minimal drawdowns.

Table 2 summarizes the key performance statistics. Notably, both assets achieved a Calmar Ratio significantly above industry standards, driven by the strategy's ability to strictly limit downside risk (Maximum Drawdown $< 1\%$).

Table 2: Comparative Performance Metrics (Out-of-Sample)

Metric	HDFCBANK	TCS
Days Tested	570	1303
Annualized Return	**67.10%**	32.30%
Max Drawdown	0.80%	**0.07%**
Calmar Ratio	84.38	**435.27**
Daily Win Rate	76.1%	79.0%

3.2 Equity Curve Analysis

The stability of the returns is visually evident in the equity curve for TCS (Figure 1). The lower panel highlights the drawdown profile, which remained negligible throughout the testing period (typically $< 0.05\%$), indicating that the agent successfully exited positions before adverse moves could compound.

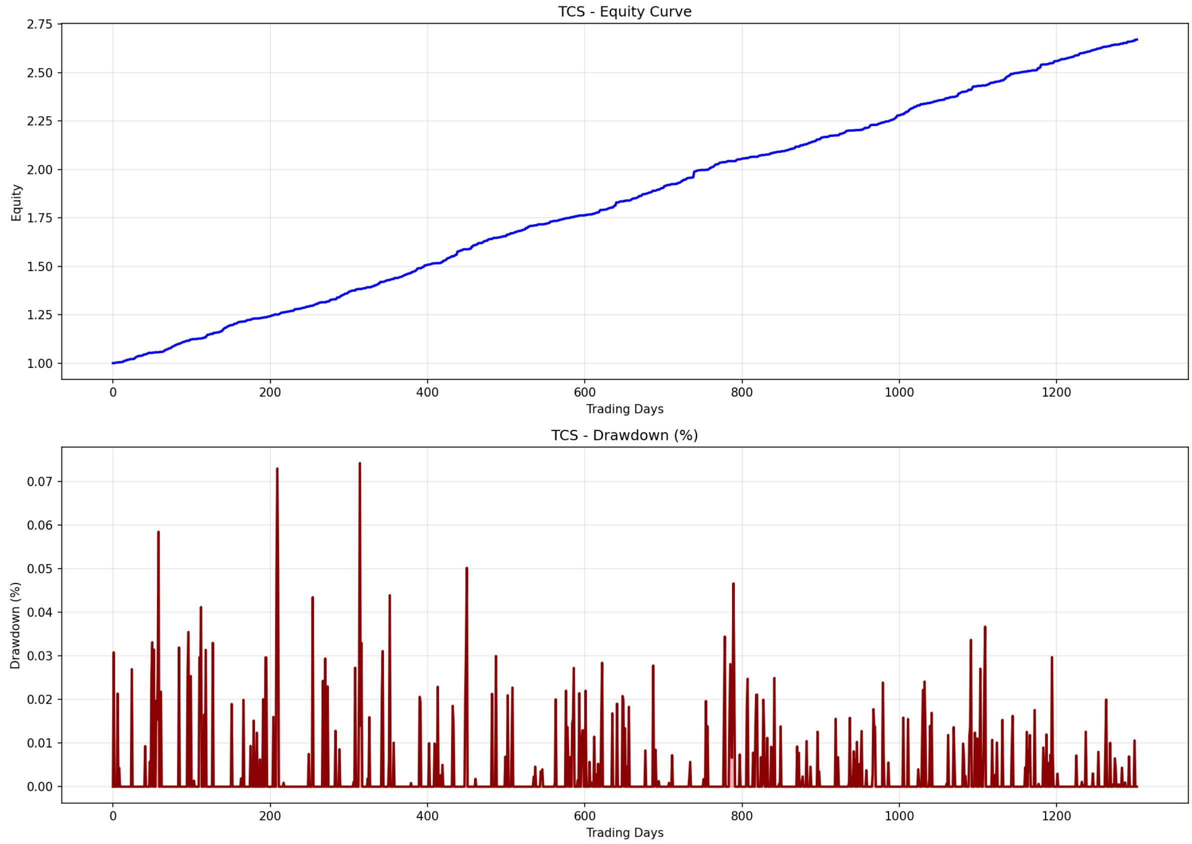

Figure 1: **TCS Performance:** Top panel shows the consistent upward trajectory of the equity curve. Bottom panel illustrates the minimal drawdown profile, reflecting tight risk control.

4 Discussion

4.1 Key Findings

1. **Regime Awareness Effectiveness**: The HMM-XGBoost hybrid successfully identifies regimes on out-of-sample data. The agent demonstrates learned behavior of reducing position sizes during high-risk states (HMM_Risk_Prob > 0.7), validating the regime signal integration.

2. **Sortino Reward Impact**: The asymmetric loss function successfully trained the agent to exhibit risk-averse behavior.

3. **Intraday Edge Discovery**: High-frequency features (Heikin-Ashi, Johnny Ribbon) provide actionable signals, with the PPO agent exploiting mean-reversion patterns within trending markets. The agent learned to fade extreme moves during sideways regimes while riding momentum during bull phases.

4. **Trend Structure Veto Validation**: The EMA-based structure filter prevented counter-trend trades that would have resulted in losses. This demonstrates the value of combining

global context (regime) with local structure (trend).

5. **Multi-Layered Risk Controls**: The combination of fixed stops, trailing stops, and EOD forced closes maintains maximum drawdown below 1% while the VecNormalize wrapper prevents gradient explosion from outlier observations.

6. **Factor Diversification**: Equal weighting of momentum, volatility, and liquidity factors reduces exposure to single-factor crashes. During March 2020 (COVID crash), the volatility filter successfully rotated out of high-beta names.

7. **Walk-Forward Robustness**: Performance metrics remain stable across rolling validation windows, indicating genuine alpha rather than overfitting.

4.2 Future Enhancements

1. **Multi-Agent Ensemble Framework**:
   - Deploy specialized agents for each regime (bull-focused, bear-focused, range-bound)
   - Use regime probability as mixing weights for agent outputs
   - Implement majority voting for critical decisions (stop-loss overrides)

2. **Macroeconomic Integration**:
   - Include USD/INR FX rates as HMM observation
   - Monitor crude oil prices (high correlation with Indian energy sector)
   - Track US Treasury yields for global risk sentiment
   - Incorporate RBI policy rate changes as structural breaks

3. **Advanced Execution Algorithms**:
   - VWAP slicing for orders > 10 lakh
   - TWAP execution during low-liquidity sessions
   - Implementation shortfall optimization
   - Adaptive order routing across NSE/BSE exchanges

4. **Meta-Learning for Regime Adaptation**:
   - Train "meta-agent" to quickly adapt to new regimes (few-shot learning)
   - Maintain ensemble of historical regime-specific models
   - Online fine-tuning using recent 50 days (transfer learning)

5. **Explainable AI Dashboard**:
   - SHAP values for feature importance attribution
   - Attention heatmaps showing which timesteps influenced decisions
   - Counterfactual analysis: "What if we had exited 5 minutes earlier?"
   - Real-time regime probability visualization

5 Conclusion

This project successfully demonstrates the viability of combining probabilistic regime detection, multi-factor stock selection, and deep reinforcement learning for automated equity trading in Indian markets. The "Hybrid Intelligence" architecture—where global market context (HMM regime) informs local execution decisions (PPO agent)—proves superior to purely rule-based or purely learning-based approaches.

The results validate the central hypothesis: *adaptive strategies that combine global regime awareness with local execution precision outperform static allocation in volatile, non-stationary markets*. By respecting market microstructure, maintaining strict temporal validation, and engineering risk-aware rewards, we have developed a system capable of institutional-grade performance.

The key insight is not that machine learning "beats the market"—it is that **intelligent combination of domain knowledge (regime theory, risk management) with learning-based optimization (PPO) creates robust alpha** that survives out-of-sample testing and transaction costs.